

# Mini-Lab 3: Process

COMP3230, The University of Hong Kong

Sept. 2024

## Total 1 point

### Objective

At the end of this mini-lab, you will be able to:

- Gain hands-on experience in process-related system calls, e.g. `fork()`, `wait()`, and `exec()`.
- Understand the parental relationship between processes.

### Instructions

In C, the `fork()` system call is used to create a new process, called the **child process**, which runs concurrently with the **parent process**. When `fork()` is called, it creates a duplicate of the parent process with a unique Process ID (PID). Both the parent and child processes execute the same code following `fork()`, but you can distinguish them using the return value: `fork()` returns 0 in the child and the child's PID in the parent.

If we do not distinguish between the child and parent processes, it may lead to potential risks such as inconsistent behavior, data duplication, or race conditions. Both processes will continue running identical code, which can result in both modifying separate copies of the same variables, leading to unpredictable outcomes or resource conflicts if shared data or external resources like files are accessed without proper coordination. For example:

```
#include <stdio.h>
#include <unistd.h>

int main() {
    int x = 10;
    fork(); // Creates a new process, duplicating the entire memory space
    x += 10; // Both processes modify their own copies of x
    printf("Value of x: %d\n", x);
    return 0;
}
```

Both parent and child modify their own copies of x, leading to potential issues if processes need to share or synchronize data, as they work on separate copies.

The `exec()` family of functions replaces the current process with a new program. Once `exec()` is called, the existing code and data in the current process are replaced, and the new program starts executing. This means none of the original code after `exec()` will run, as the process memory is entirely replaced with the new program. The process keeps the same PID, but everything else is swapped out.

```
#include <unistd.h>

int main() {
    execl("/bin/ls", "ls", NULL);
    printf("This will not be printed\n"); // This line will never be
                                         executed
    return 0;
}
```

In this example, after `execl()` is called to run the ls command, the rest of the code (like the `printf()` line) is not executed.

In this mini-lab, you will implement a simple multiprocessing program and `fork()` and `exec()` to understand how the operating system manages multiple processes.

1. Open two terminals marked as Terminal 1 and Terminal 2.
2. In Terminal 1, find out and record *process id* of the current shell using  
`echo $$ > t1-pid.txt`.
3. Complete all TODO sections in `lab3-process.c` using `fork()`, `exec()`, and `wait()` following the inline instructions.
4. In Terminal 1, compile and execute `lab3-process.c`. If your implementation is correct, the child process will sleep 20 seconds.
5. Before the process returns in Terminal 1, in Terminal 2, execute `pstree -sp $(cat t1-pid.txt)` to monitor the tree of running processes. We can see a chain of processes like SSH > Bash (Terminal 1) > parent process (`lab3-process`) > child process (`sleep 20`).

**Explanation of their relations:** Upon connecting remotely, the SSH process forks to create a new Bash shell for the user's session. When `lab3-process` is executed within this Bash shell, Bash forks again to run the program, which `lab3-process` forks once more to create a child process for the `sleep 20` command. This sequence of fork operations establishes the chain: SSH forks Bash > Bash forks Parent (`lab3-process`) > Parent forks Child (`sleep 20`).

## **Submission**

(1 pt) Complete all TODO sections and submit your code as `lab3-process_<your_student_id>.c`.

## Appendix

```
// file: lab3-process.c
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t pid = -1;
    printf("Before fork(). variable pid = %d, getpid()=%d, getppid()=%d\n",
           pid, getpid(), getppid());
    fflush(stdout);
    // TODO: Create a child process using fork() and store the return value
    // in pid (~1 line)

    if (pid < 0) {
        fprintf(stderr, "fork() Failed");
        exit(-1);
    } else if (pid == 0) {
        // Child Process
        printf("Child Process. variable pid = %d, getpid()=%d, getppid()=%d
               \n", pid, getpid(), getppid());
        // TODO: Use exec() family to replace the current process image
        // with "sleep", "20" (~1 line)

    } else {
        // TODO: Make the parent process wait for the child to complete (~1
        // line)

        printf("Parent Process. variable pid = %d, getpid()=%d, getppid()=%d
               \n", pid, getpid(), getppid());
    }
    return 0;
} // main
```